

DSDE-Homework-1_Huggingface_image_classification

fine-tune ViT สำหรับทำ classification input เป็นใบของต้นถั่ว beans data set ด้วย Hugging Face library โหลด datasets

```
from datasets import load_dataset
ds = load_dataset('beans')
ds
```

```
image = ds['train'][400]['image']
image
```



function สำหรับดู data ใน dataset

```
from transformers.utils.dummy_vision_objects import ImageGPTFeatureExtractor
import random
from PIL import ImageDraw, ImageFont, Image

def show_examples(ds, seed: int = 1234, examples_per_class: int = 3, size=(350, 350)):

    w, h = size
    labels = ds['train'].features['labels'].names
    grid = Image.new('RGB', size=(examples_per_class * w, len(labels) * h))
    draw = ImageDraw.Draw(grid)
    font = ImageFont.truetype("/usr/share/fonts/truetype/liberation/LiberationMono-Bold.ttf", 24)

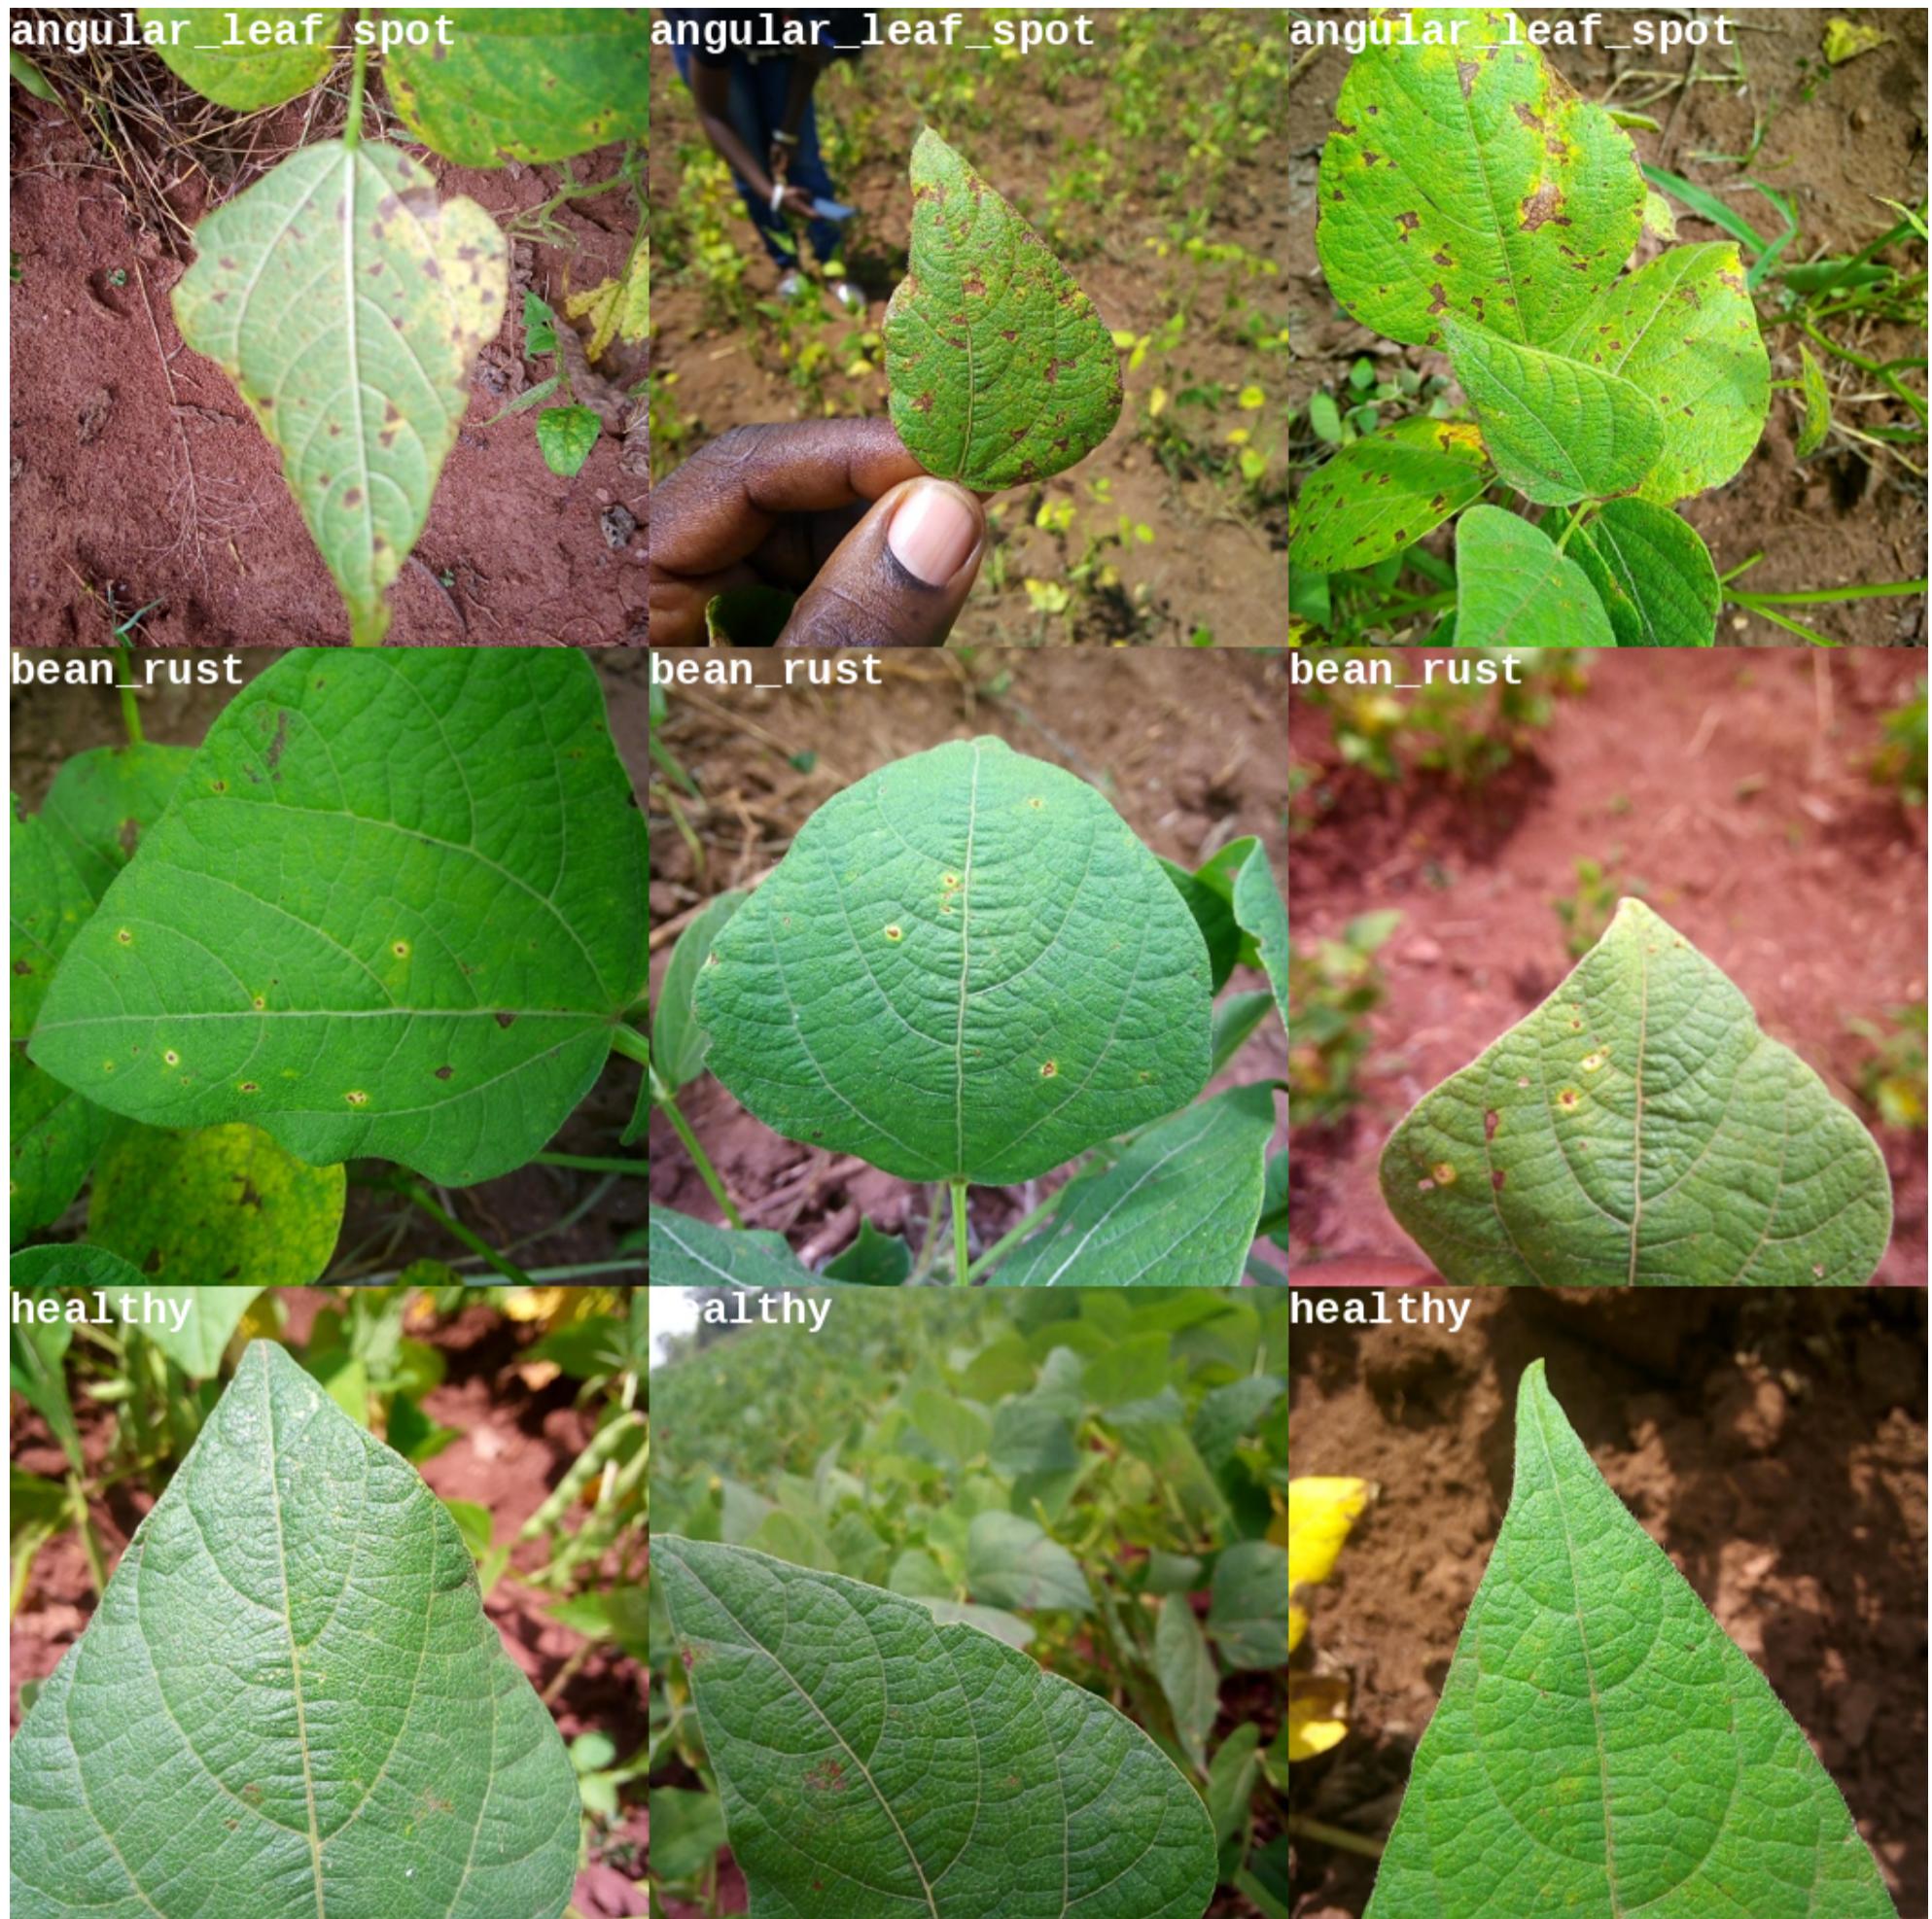
    for label_id, label in enumerate(labels):

        # Filter the dataset by a single label, shuffle it, and grab a few samples
        ds_slice = ds['train'].filter(lambda ex: ex['labels'] ==
label_id).shuffle(seed).select(range(examples_per_class))

        # Plot this label's examples along a row
        for i, example in enumerate(ds_slice):
            image = example['image']
            idx = examples_per_class * label_id + i
            box = (idx % examples_per_class * w, idx // examples_per_class * h)
            grid.paste(image.resize(size), box=box)
            draw.text(box, label, (255, 255, 255), font=font)

    return grid
```

```
show_examples(ds, seed=random.randint(0, 1337), examples_per_class=3)
```



ໂຫລດ pretrained model

```
from transformers import ViTImageProcessor
model_name_or_path = 'google/vit-base-patch16-224'
feature_extractor = ViTImageProcessor.from_pretrained(model_name_or_path)
feature_extractor
```

```
ViTImageProcessor {
    "do_normalize": true,
    "do_rescale": true,
    "do_resize": true,
    "image_mean": [
        0.5,
        0.5,
        0.5
    ],
    "image_processor_type": "ViTImageProcessor",
```

```
"image_std": [
    0.5,
    0.5,
    0.5
],
"resample": 2,
"rescale_factor": 0.00392156862745098,
"size": {
    "height": 224,
    "width": 224
}
}
```

function สำหรับ transform

```
def transform(example_batch):
    # Take a list of PIL images and turn them to pixel values
    inputs = feature_extractor([x for x in example_batch['image']], return_tensors='pt')

    # Don't forget to include the labels!
    inputs['labels'] = example_batch['labels']

    return inputs

prepared_ds = ds.with_transform(transform)
```

```
import torch
```

```
def collate_fn(batch):
    return { 'pixel_values': torch.stack([x['pixel_values'] for x in batch]),
        'labels': torch.tensor([x['labels'] for x in batch])}
```

load accuracy metric สำหรับ evaluation

```
import numpy as np
from datasets import load_metric

metric = load_metric("accuracy")

def compute_metrics(p):
    return metric.compute(predictions=np.argmax(p.predictions, axis=1), references=p.label_ids)
```

โหลด pretrain สำหรับ classification

```
from transformers import ViTForImageClassification

labels = ds['train'].features['labels'].names
model = ViTForImageClassification.from_pretrained('google/vit-base-patch16-224',
    num_labels=len(labels),
```

```
ignore_mismatched_sizes=True)
```

ประการ Argument ການ train

```
from transformers import TrainingArguments

training_args = TrainingArguments(
    output_dir="vit-base-beans-demo-v5",
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    evaluation_strategy="epoch",
    save_strategy="epoch",
    logging_strategy='epoch',
    logging_steps=1,
    num_train_epochs=10,
    optim='adamw_torch',
    learning_rate=1e-5,
    remove_unused_columns=False,
    report_to='tensorboard',
    load_best_model_at_end=True,
    metric_for_best_model="loss",
    seed=0
)
```

ประการ trainer

```
from transformers import Trainer

trainer = Trainer(
    model=model,
    args=training_args,
    data_collator=collate_fn,
    compute_metrics=compute_metrics,
    train_dataset=prepared_ds["train"],
    eval_dataset=prepared_ds["validation"],
    tokenizer=feature_extractor,
)
```

ທດລວງ evaluate กອນ fine tune

```
from sklearn.metrics import classification_report

test_predictions = trainer.predict(prepared_ds['test'])

# For each prediction, create the label with argmax

predictions = np.argmax(test_predictions[0], axis=1)

# Retrieve reference labels from test set

test_labels = np.array(prepared_ds["test"][:, ["labels"]])

print(classification_report(test_labels, predictions, target_names=['angular_leaf_spot', 'bean_rust', 'healthy'],
digits=4))
```

	precision	recall	f1-score	support
angular_leaf_spot	0.0000	0.0000	0.0000	43
bean_rust	0.3097	0.8140	0.4487	43
healthy	0.3333	0.0714	0.1176	42
accuracy			0.2969	128
macro avg	0.2144	0.2951	0.1888	128
weighted avg	0.2134	0.2969	0.1893	128

save result

```
train_results = trainer.train()

trainer.save_model()

trainer.log_metrics("train", train_results.metrics)

trainer.save_metrics("train", train_results.metrics)

trainer.save_state()
```

```
metrics = trainer.evaluate(prepared_ds['validation'])
trainer.log_metrics("eval", metrics)
trainer.save_metrics("eval", metrics)
```

***** eval metrics *****

epoch	=	10.0
eval_accuracy	=	0.9925
eval_loss	=	0.0242
eval_runtime	=	0:00:03.01
eval_samples_per_second	=	44.081
eval_steps_per_second	=	2.983

prediction

```
from sklearn.metrics import classification_report

test_predictions = trainer.predict(prepared_ds['test'])
# For each prediction, create the label with argmax
predictions = np.argmax(test_predictions[0], axis=1)
# Retrieve reference labels from test set
test_labels = np.array(prepared_ds["test"][:, ["labels"]])
print(classification_report(test_labels, predictions, target_names=['angular_leaf_spot', 'bean_rust', 'healthy'],
digits=4))
```

tensorboard

```
%load_ext tensorboard
%tensorboard --logdir '{"vit-base-beans-demo-v5"}' /runs
```

